                    Diffie-Hellman Group Exchange for
              the Secure Shell (SSH) Transport Layer Protocol

Status of This Memo

   This document specifies an Internet standards track protocol for the
   Internet community, and requests discussion and suggestions for
   improvements. Please refer to the current edition of the "Internet
   Official Protocol Standards" (STD 1) for the standardization state
   and status of this protocol. Distribution of this memo is unlimited.

Copyright Notice

Abstract

   This memo describes a new key exchange method for the Secure Shell
   (SSH) protocol. It allows the SSH server to propose new groups on
   which to perform the Diffie-Hellman key exchange to the client. The
   proposed groups need not be fixed and can change with time.

1. Overview and Rationale

   SSH [RFC4251] is a very common protocol for secure remote login on
   the Internet. Currently, SSH performs the initial key exchange using
   the "diffie-hellman-group1-sha1" method [RFC4253]. This method
   prescribes a fixed group on which all operations are performed.

   The Diffie-Hellman key exchange provides a shared secret that cannot
   be determined by either party alone. Furthermore, the shared secret
   is known only to the participant parties. In SSH, the key exchange
   is signed with the host key to provide host authentication.

   The security of the Diffie-Hellman key exchange is based on the
   difficulty of solving the Discrete Logarithm Problem (DLP). Since we
   expect that the SSH protocol will be in use for many years in the
   future, we fear that extensive precomputation and more efficient
   algorithms to compute the discrete logarithm over a fixed group might
   pose a security threat to the SSH protocol.

The ability to propose new groups will reduce the incentive to use
precomputation for more efficient calculation of the discrete
logarithm. The server can constantly compute new groups in the
background.

2. Requirements Notation

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
"SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
document are to be interpreted as described in [RFC2119].

3. Diffie-Hellman Group and Key Exchange

The server keeps a list of safe primes and corresponding generators
that it can select from. A prime p is safe if p = 2q + 1 and q is
prime. New primes can be generated in the background.

The generator g should be chosen such that the order of the generated
subgroup does not factor into small primes; that is, with p = 2q + 1,
the order has to be either q or p - 1. If the order is p - 1, then
the exponents generate all possible public values, evenly distributed
throughout the range of the modulus p, without cycling through a
smaller subset. Such a generator is called a "primitive root" (which
is trivial to find when p is "safe").

The client requests a modulus from the server indicating the
preferred size. In the following description (C is the client, S is
the server, the modulus p is a large safe prime, and g is a generator
for a subgroup of GF(p), min is the minimal size of p in bits that is
acceptable to the client, n is the size of the modulus p in bits that
the client would like to receive from the server, max is the maximal
size of p in bits that the client can accept, V_S is S's version
string, V_C is C's version string, K_S is S's public host key, I_C is
C's KEXINIT message, and I_S is S's KEXINIT message that has been
exchanged before this part begins):

1. C sends "min || n || max" to S, indicating the minimal acceptable
    group size, the preferred size of the group, and the maximal
    group size in bits the client will accept.

2. S finds a group that best matches the client's request, and sends
    "p || g" to C.

3. C generates a random number x, where 1 < x < (p-1)/2. It
    computes e = g^x mod p, and sends "e" to S.

   4. S generates a random number y, where 0 < y < (p-1)/2, and
       computes f = g^y mod p. S receives "e". It computes K = e^y mod
       p, H = hash(V_C || V_S || I_C || I_S || K_S || min || n || max ||
       p || g || e || f || K) (these elements are encoded according to
       their types; see below), and signature s on H with its private
       host key. S sends "K_S || f || s" to C. The signing operation
       may involve a second hashing operation.

   5. C verifies that K_S really is the host key for S (e.g., using
       certificates or a local database to obtain the public key). C is
       also allowed to accept the key without verification; however,
       doing so will render the protocol insecure against active attacks
       (but may be desirable for practical reasons in the short term in
       many environments). C then computes K = f^x mod p, H = hash(V_C
       || V_S || I_C || I_S || K_S || min || n || max || p || g || e ||
       f || K), and verifies the signature s on H.

   Servers and clients SHOULD support groups with a modulus length of k
   bits, where 1024 <= k <= 8192. The recommended values for min and
   max are 1024 and 8192, respectively.

   Either side MUST NOT send or accept e or f values that are not in the
   range [1, p-1]. If this condition is violated, the key exchange
   fails. To prevent confinement attacks, they MUST accept the shared
   secret K only if 1 < K < p - 1.

   The server should return the smallest group it knows that is larger
   than the size the client requested. If the server does not know a
   group that is larger than the client request, then it SHOULD return
   the largest group it knows. In all cases, the size of the returned
   group SHOULD be at least 1024 bits.

   This is implemented with the following messages. The hash algorithm
   for computing the exchange hash is defined by the method name, and is
   called HASH. The public key algorithm for signing is negotiated with
   the KEXINIT messages.

   First, the client sends:

     byte   SSH_MSG_KEY_DH_GEX_REQUEST
     uint32 min, minimal size in bits of an acceptable group
     uint32 n, preferred size in bits of the group the server will send
     uint32 max, maximal size in bits of an acceptable group

The server responds with

```
byte    SSH_MSG_KEX_DH_GEX_GROUP
mpint   p, safe prime
mpint   g, generator for subgroup in GF(p)
```

The client responds with:

```
byte    SSH_MSG_KEX_DH_GEX_INIT
mpint   e
```

The server responds with:

```
byte    SSH_MSG_KEX_DH_GEX_REPLY
string  server public host key and certificates (K_S)
mpint   f
string  signature of H
```

The hash H is computed as the HASH hash of the concatenation of the following:

```
string  V_C, the client's version string (CR and NL excluded)
string  V_S, the server's version string (CR and NL excluded)
string  I_C, the payload of the client's SSH_MSG_KEXINIT
string  I_S, the payload of the server's SSH_MSG_KEXINIT
string  K_S, the host key
uint32  min, minimal size in bits of an acceptable group
uint32  n, preferred size in bits of the group the server will send
uint32  max, maximal size in bits of an acceptable group
mpint   p, safe prime
mpint   g, generator for subgroup
mpint   e, exchange value sent by the client
mpint   f, exchange value sent by the server
mpint   K, the shared secret
```

This value is called the exchange hash, and it is used to authenticate the key exchange as per [RFC4253].

4. Key Exchange Methods

This document defines two new key exchange methods:
"diffie-hellman-group-exchange-sha1" and
"diffie-hellman-group-exchange-sha256".

4.1. diffie-hellman-group-exchange-sha1

   The "diffie-hellman-group-exchange-sha1" method specifies
   Diffie-Hellman Group and Key Exchange with SHA-1 [FIPS-180-2] as
   HASH.

4.2. diffie-hellman-group-exchange-sha256

   The "diffie-hellman-group-exchange-sha256" method specifies
   Diffie-Hellman Group and Key Exchange with SHA-256 [FIPS-180-2] as
   HASH.

   Note that the hash used in key exchange (in this case, SHA-256) must
   also be used in the key derivation pseudo-random function (PRF), as
   per the requirement in the "Output from Key Exchange" section in
   [RFC4253].

5. Summary of Message Numbers

   The following message numbers have been defined in this document.
   They are in a name space private to this document and not assigned by
   IANA.

```
   #define SSH_MSG_KEX_DH_GEX_REQUEST_OLD  30
   #define SSH_MSG_KEX_DH_GEX_REQUEST      34
   #define SSH_MSG_KEX_DH_GEX_GROUP        31
   #define SSH_MSG_KEX_DH_GEX_INIT         32
   #define SSH_MSG_KEX_DH_GEX_REPLY        33
```

   SSH_MSG_KEX_DH_GEX_REQUEST_OLD is used for backward compatibility.
   Instead of sending "min || n || max", the client only sends "n". In
   addition, the hash is calculated using only "n" instead of "min || n
   || max".

   The numbers 30-49 are key exchange specific and may be redefined by
   other kex methods.

6. Implementation Notes

6.1. Choice of Generator

   One useful technique is to select the generator, and then limit the
   modulus selection sieve to primes with that generator:

```
   2   when p (mod 24) = 11.
   5   when p (mod 10) = 3 or 7.
```

It is recommended to use 2 as generator, because it improves
efficiency in multiplication performance. It is usable even when it
is not a primitive root, as it still covers half of the space of
possible residues.

6.2. Private Exponents

To increase the speed of the key exchange, both client and server may
reduce the size of their private exponents. It should be at least
twice as long as the key material that is generated from the shared
secret. For more details, see the paper by van Oorschot and Wiener
[VAN-OORSCHOT].

7. Security Considerations

This protocol aims to be simple and uses only well-understood
primitives. This encourages acceptance by the community and allows
for ease of implementation, which hopefully leads to a more secure
system.

The use of multiple moduli inhibits a determined attacker from
precalculating moduli exchange values, and discourages dedication of
resources for analysis of any particular modulus.

It is important to employ only safe primes as moduli, as they allow
us to find a generator g so that the order of the generated subgroup
does not factor into small primes, that is, with p = 2q + 1, the
order has to be either q or p - 1. If the order is p - 1, then the
exponents generate all possible public values, evenly distributed
throughout the range of the modulus p, without cycling through a
smaller subset. Van Oorshot and Wiener note that using short private
exponents with a random prime modulus p makes the computation of the
discrete logarithm easy [VAN-OORSCHOT]. However, they also state
that this problem does not apply to safe primes.

The least significant bit of the private exponent can be recovered
when the modulus is a safe prime [MENEZES]. However, this is not a
problem if the size of the private exponent is big enough. Related
to this, Waldvogel and Massey note: When private exponents are chosen
independently and uniformly at random from {0,...,p-2}, the key
entropy is less than 2 bits away from the maximum, lg(p-1)
[WALDVOGEL].

The security considerations in [RFC4251] also apply to this key
exchange method.

8. Acknowledgements

   The document is derived in part from "SSH Transport Layer Protocol"
   [RFC4253] by T. Ylonen, T. Kivinen, M. Saarinen, T. Rinne, and S.
   Lehtinen.

   Markku-Juhani Saarinen pointed out that the least significant bit of
   the private exponent can be recovered efficiently when using safe
   primes and a subgroup with an order divisible by two.

   Bodo Moeller suggested that the server send only one group, reducing
   the complexity of the implementation and the amount of data that
   needs to be exchanged between client and server.

Appendix A: Generation of Safe Primes

   The "Handbook of Applied Cryptography" [MENEZES] lists the following
   algorithm to generate a k-bit safe prime p. It has been modified so
   that 2 is a generator for the multiplicative group mod p.

   1. Do the following:

       1. Select a random (k-1)-bit prime q, so that q mod 12 = 5.

       2. Compute p := 2q + 1, and test whether p is prime (using,
           e.g., trial division and the Rabin-Miller test).

   2. Repeat until p is prime.

   If an implementation uses the OpenSSL libraries, a group consisting
   of a 1024-bit safe prime and 2 as generator can be created as
   follows:

       DH *d = NULL;
       d = DH_generate_parameters(1024, DH_GENERATOR_2, NULL, NULL);
       BN_print_fp(stdout, d->p);

   The order of the subgroup generated by 2 is q = p - 1.

References

Normative References

   [FIPS-180-2]   National Institute of Standards and Technology (NIST),
                  "Secure Hash Standard (SHS)", FIPS PUB 180-2,
                  August 2002.

   [RFC4251]   Ylonen, T. and C. Lonvick, "The Secure Shell (SSH)
                  Protocol Architecture", RFC 4251, January 2006.

   [RFC4253]   Lonvick, C., "The Secure Shell (SSH) Transport Layer
                  Protocol", RFC 4253, January 2006.

   [RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate
                  Requirement Levels", BCP 14, RFC 2119, March 1997.

Informative References

   [MENEZES]   Menezes, A., van Oorschot, P., and S. Vanstone,
                  "Handbook of Applied Cryptography", CRC Press, p. 537,
                  1996.

   [VAN-OORSCHOT] van Oorschot, P. and M. Wiener, "On Diffie-Hellman key
                   agreement with short exponents", Advances in
                   Cryptology -EUROCRYPT'96, LNCS 1070,
                   Springer-Verlag, pp. 332-343, 1996.

   [WALDVOGEL]   Waldvogel, C. and J. Massey, "The probability
                   distribution of the Diffie-Hellman key", Proceedings
                   of AUSCRYPT 92, LNCS 718, Springer-Verlag, pp.
                   492-504, 1993.

Authors' Addresses

   Markus Friedl
   EMail: markus@openbsd.org

   Niels Provos
   EMail: provos@citi.umich.edu

   William A. Simpson
   EMail: wsimpson@greendragon.com